

CSc 360: Operating Systems (Fall 2019)

Programming Assignment 3 (P3): A Simple File System (SFS)

Spec out: Nov 1, 2019
Due Date: Nov 29, 2019

1 Introduction

So far, you have built a shell environment and a multi-thread scheduler with process synchronization. Excellent job! What is still missing for a “real” operating system? A file system! In this assignment, you will implement utilities that perform operations on a file system similar to Microsoft’s FAT file system with some improvement.

1.1 Sample File Systems

You will be given a test file system disk image for self-testing, but you can create your own image following the specification, and your submission may be tested against other disk images following the same specification.

You should get comfortable examining the raw, **binary** data in the file system images using the program `xxd`.

VERY IMPORTANT: since you are dealing with binary data, functions intended for string manipulation such as `strcpy()` do NOT work (since binary data may contain binary ‘0’ anywhere), and you should use functions intended for binary data such as `memcpy()`.

2 Tutorial Schedule

In order to help you finish this programming assignment on time successfully, the schedule of the lectures and the tutorials has been adjusted. There are two tutorials and one help session arranged during the course of this assignment. **NOTE: Please do attend the tutorials and follow the tutorial schedule closely.**

Date	Tutorial	Milestones
Nov 5/7/8	P3 spec and practice questions go-through	design done
Nov 12/14/15	reading break, no tutorials	coding half done
Nov 19/21/22	more on design and implementation	coding done
Nov 26/28/29	no tutorials but a help session	final deliverable

3 Requirements

3.1 Part I (3 points)

In part I, you will write a program that displays information about the file system. In order to complete part I, you will need to read the file system super block and use the information in the super block to read the FAT.

Your program for part I will be invoked as follows (output value here just for illustration purposes):

```
./diskinfo test.img
```

Sample output:

```
Super block information:
Block size: 512
Block count: 5120
FAT starts: 1
```

```
FAT blocks: 40
Root directory start: 41
Root directory blocks: 8
```

```
FAT information:
Free Blocks: 5071
Reserved Blocks: 41
Allocated Blocks: 8
```

Please be sure to use the exact same output **format** as shown above.

3.2 Part II (3 points)

In part II, you will write a program, with the routines already implemented for part I, that displays the contents of the root directory or a given sub-directory in the file system.

Your program for part II will be invoked as follows:

```
./disklist test.img /sub_dir
```

The directory listing should be **formatted** as follows:

1. The first column will contain:

(a) **F** for regular files, or

(b) **D** for directories;

followed by a single space

2. then 10 characters to show the file size, followed by a single space

3. then 30 characters for the file name, followed by a single space

4. then the file modification date (we will not display the file creation date).

For example:

F	2560	foo.txt	2005/11/15 12:00:00
F	5120	foo2.txt	2005/11/15 12:00:00
F	48127	makefs	2005/11/15 12:00:00
F	8	foo3.txt	2005/11/15 12:00:00

3.3 Part III (3 points)

In part III, you will write a program that copies a file from the file system to the current directory in Linux. If the specified file is not found in the root directory or a given sub-directory of the file system, you should output the message **File not found.** and exit.

Your program for part III will be invoked as follows:

```
./diskget test.img /sub_dir/foo2.txt foo.txt
```

3.4 Part IV (3 points)

In part IV, you will write a program that copies a file from the current Linux directory into the file system, at the root directory or a given sub-directory. If the specified file is not found, you should output the message **File not found.** on a single line and exit.

Your program for part IV will be invoked as follows:

```
./diskput test.img foo.txt /sub_dir/foo3.txt
```

3.5 Part V (3 points)

From time to time, the disk image may get corrupted due to wrong or incomplete operations, including what happened to the test image, although it does not affect the first four parts. Your fifth part is to go through the disk image according to the file system specification, including the super block, FDT, FAT and data blocks, find inconsistent information among them and fix these issues when possible. For example, a block indicated as reserved (for FAT and root directory) in FAT might be mistakenly used as a data block. In this case, the data shall be relocated, and the FAT and possibly FDT are updated accordingly. Also, a block indicated as allocated in FAT does not belong to any files. In this case, the entry in FAT is fixed to available. Further, a block is the last block of a file according to its size, but it is not indicated by -1 in FAT. In this case, the FAT entry is updated to -1. On the other hand, a block is not the last block of a file according to its size, but it is indicated by -1 in FAT, so the file is truncated up to this last block. Your program needs to be able to handle at least these three cases.

Your program for part V will be invoked as follows:

```
./diskfix test.img
```

and output the problems that you identified and possibly fixed, e.g.,

```
Block 5 indicated reserved in FAT but used by foo.txt; foo.txt relocated
Block 1005 indicated allocated in FAT but not used by any files; fixed to available
Block 2005 is the last block of foo2.txt but not indicated -1 in FAT; fixed to -1
Block 3005 is not the last block of foo3.txt but indicated -1 in FAT; foo3.txt truncated to 4096 bytes
```

4 File System Specification

The FAT file system has three major components:

1. the super block,
2. the directory structure.
3. the File Allocation Table (informally referred to as the FAT),

Each of these three components is described in the subsections below.

4.1 File System Superblock

The first block (512 bytes) is reserved to contain information about the file system. The layout of the superblock is as follows:

Description	Size	Default Value
File system identifier	8 bytes	CSC360FS
Block Size	2 bytes	0x200
File system size (in blocks)	4 bytes	0x00001400
Block where FAT starts	4 bytes	0x00000001
Number of blocks in FAT	4 bytes	0x00000028
Block where root directory starts	4 bytes	0x00000029
Number of blocks in root dir	4 bytes	0x00000008

Figure 1: Superblock Fields

Note: Block number starts from 0 in the file system.

4.2 Directory Entries

Each directory entry takes up 64 bytes, which implies there are 8 directory entries per 512 byte block.

Each directory entry has the following structure:

The description of each field follows:

Description	Size
Status	1 byte
Starting Block	4 bytes
Number of Blocks	4 bytes
File Size (in bytes)	4 bytes
Create Time	7 bytes
Modify Time	7 bytes
File Name	31 bytes
unused (set to 0xFF)	6 bytes

Figure 2: Directory Entry

Bit 0	set to 0 if this directory entry is available, set to 1 if it is in use
Bit 1	set to 1 if this entry is a normal file
Bit 2	set to 1 if this entry is a directory

Figure 3: Format of Status Field

Status This is bit mask that is used to describe the status of the file. Currently only 3 of the bits are used.

It is implied that only one of bit 2 or bit 1 can be set to 1. That is, an entry is either a normal file or it is a directory, *not both*.

Starting Block This is the location on disk of the first block in the file

Number of Blocks The total number of blocks in this file

File Size The size of the file, in bytes. The size of this field implies that the largest file we can support is 2^{32} bytes long.

Create Time The date and time when this file was created. The file system stores the system times as integer values in the format:

YYYYMMDDHHMMSS

Field	Size
YYYY	2 bytes
MM	1 byte
DD	1 byte
HH	1 byte
MM	1 byte
SS	1 byte

Figure 4: Format of Date-Time Field

Modify Time The last time this file was modified. Stored in the same format as the Create Time shown above.

File Name The file name, null terminated. Because of the null terminator, the maximum length of any filename is 30 bytes.

Valid characters are upper and lower case letters (a-z, A-Z), digits (0-9) and the underscore character (-).

4.3 File Allocation Table (FAT)

Each directory entry contains the starting block number for a file, let's say it is block number X. To find the next block in the file, you should look at entry X in the FAT. If the value you find there does not indicate End-of-File (see below) then that value, call it Y, is the next block number in the file.

That is, the first block is at block number X, you look in the FAT table at entry X and find the value Y. The second data block is at block number Y. Then you look in the FAT at entry Y to find the next data block number... continue this until you find the special value in the FAT entry indicating that you are at the last FAT entry of the file.

The FAT is really just a linked list, which the head of the list being stored in the "Starting Block" field in the directory entry, and the 'next pointers' being stored in the FAT entries.

FAT entries are 4 bytes long (32 bits), which implies there are 128 FAT entries per block.
Special values for FAT entries are described in Figure 5.

Value	Meaning
0x00000000	This block is available
0x00000001	This block is reserved
0x00000002– 0xFFFFFFFF00	Allocated blocks as part of files
0xFFFFFFFF	This is the last block in a file

Figure 5: Value of FAT entry

5 Byte Ordering

Different hardware architectures store multi-byte data (like integers) in different orders. Consider the large integer: 0xDEADBEEF

On the Intel architecture (Little Endian), it would be stored in memory as:

EF BE AD DE

On the PowerPC (Big Endian), it would be stored in memory as:

DE AD BE EF

Our file system will use Big Endian for storage. This will make debugging the file system by examining the raw data much easier.

This will mean that you have to convert all your integer values to Big Endian before writing them to disk. There are utility functions in `netinit/in.h` that do exactly that. (When sending data over the network, it is expected the data is in Big Endian format.)

See the functions `htons`, `htonl`, `ntohs` and `ntohl`.

The side effect of using these functions will be that your code will work on multiple platforms. (On machines that natively store integers in Big Endian format, like the Mac (not the Intel-based ones), the above functions don't actually do anything but you should still use them!)

6 Submission Requirements

What to hand in: You need to hand in a `.tar.gz` file containing all your source code and a Makefile that produces the executables for parts I – V.

Please include a `readme.txt` file that explains your design and implementation.

The file is submitted through `connex.csc.uvic.ca` site.

119 A An Exercise

120 Q1 Consider the superblock shown below:

```
121 0000000: 4353 4333 3630 4653 0200 0000 1400 0000 CSC360FS.....
122 0000010: 0001 0000 0028 0000 0029 0000 0008 0000 .....(....).....
123 0000020: 0000 0000 0000 0000 0000 0000 0000 0000 .....
```

124 (a) What block does the FAT start on? How many blocks are used for the FAT?

125 (b) What block does the root directory start on? How many blocks are used for the root directory?

126 Q2 Consider the following block from the root directory:

```
127      0005200: 0300 0000 3100 0000 0500 000a 0007 d50b ....1.....
128      0005210: 0f0c 0000 07d5 0b0f 0c00 0066 6f6f 2e74 .....foo.t
129      0005220: 7874 0000 0000 0000 0000 0000 0000 0000 xt.....
130      0005230: 0000 0000 0000 0000 0000 00ff ffff ffff .....
131      0005240: 0300 0000 3600 0000 0a00 0014 0007 d50b ....6.....
132      0005250: 0f0c 0000 07d5 0b0f 0c00 0066 6f6f 322e .....foo2.
133      0005260: 7478 7400 0000 0000 0000 0000 0000 0000 txt.....
134      0005270: 0000 0000 0000 0000 0000 00ff ffff ffff .....
135      0005280: 0300 0000 4000 0000 5e00 00bb ff07 d50b ....@...^.....
136      0005290: 0f0c 0000 07d5 0b0f 0c00 006d 616b 6566 .....makef
137      00052a0: 7300 0000 0000 0000 0000 0000 0000 0000 s.....
138      00052b0: 0000 0000 0000 0000 0000 00ff ffff ffff .....
139      00052c0: 0300 0000 9e00 0000 0100 0000 0807 d50b .....
140      00052d0: 0f0c 0000 07d5 0b0f 0c00 0066 6f6f 332e .....foo3.
141      00052e0: 7478 7400 0000 0000 0000 0000 0000 0000 txt.....
142      00052f0: 0000 0000 0000 0000 0000 00ff ffff ffff .....
```

143 (a) How many files are allocated in this directory? What are their names?

144 (b) How many blocks does the file makefs occupy on the disk?

145 Q3 Given the root directory information from the previous question and the FAT table shown below:

```

146      0000200: 0000 0001 0000 0001 0000 0001 0000 0001 .....
147      0000210: 0000 0001 0000 0001 0000 0001 0000 0001 .....
148      0000220: 0000 0001 0000 0001 0000 0001 0000 0001 .....
149      0000230: 0000 0001 0000 0001 0000 0001 0000 0001 .....
150      0000240: 0000 0001 0000 0001 0000 0001 0000 0001 .....
151      0000250: 0000 0001 0000 0001 0000 0001 0000 0001 .....
152      0000260: 0000 0001 0000 0001 0000 0001 0000 0001 .....
153      0000270: 0000 0001 0000 0001 0000 0001 0000 0001 .....
154      0000280: 0000 0001 0000 0001 0000 0001 0000 0001 .....
155      0000290: 0000 0001 0000 0001 0000 0001 0000 0001 .....
156      00002a0: 0000 0001 0000 002a 0000 002b 0000 002c .....*...+,
157      00002b0: 0000 002d 0000 002e 0000 002f 0000 0030 ...-...../...0
158      00002c0: ffff ffff 0000 0032 0000 0033 0000 0034 .....2...3...4
159      00002d0: 0000 0035 ffff ffff 0000 0037 0000 0038 ...5.....7...8
160      00002e0: 0000 0039 0000 003a 0000 003b 0000 003c ...9...:...;<
161      00002f0: 0000 003d 0000 003e 0000 003f ffff ffff ...=...>...?....

```

- 162 (a) What blocks does the file `foo.txt` occupy on the disk?
- 163 (b) What blocks does the file `foo2.txt` occupy on the disk?